

R vs. Julia in MCMC

Lijun Wang

29 August, 2018

Let me use the Example 7.1.1 of [Robert and Casella \(2013\)](#) to compare the speed of R and Julia in MCMC.

Example 7.1.1 Bivariate Gibbs sampler. Let the random variables X and Y have joint density $f(x, y)$, and generate a sequence of observations according to the following:

Set $X_0 = x_0$, and for $t = 1, 2, \dots$, generate

$$(7.1.1) \quad \begin{aligned} Y_t &\sim f_{Y|X}(\cdot|x_{t-1}), \\ X_t &\sim f_{X|Y}(\cdot|y_t), \end{aligned}$$

where $f_{Y|X}$ and $f_{X|Y}$ are the conditional distributions. The sequence (X_t, Y_t) , is a Markov chain, as is each sequence (X_t) and (Y_t) individually. For example, the chain (X_t) has transition density

$$K(x, x^*) = \int f_{Y|X}(y|x) f_{X|Y}(x^*|y) dy,$$

with invariant density $f_X(\cdot)$. (Note the similarity to Eaton's transition (4.10.9).)

For the special case of the bivariate normal density,

$$(7.1.2) \quad (X, Y) \sim \mathcal{N}_2 \left(0, \begin{pmatrix} 1 & \rho \\ \rho & 1 \end{pmatrix} \right),$$

the Gibbs sampler is

Given y_t , generate

$$(7.1.3) \quad \begin{aligned} X_{t+1} | y_t &\sim \mathcal{N}(\rho y_t, 1 - \rho^2), \\ Y_{t+1} | x_{t+1} &\sim \mathcal{N}(\rho x_{t+1}, 1 - \rho^2). \end{aligned}$$

The Gibbs sampler is obviously not necessary in this particular case, as iid copies of (X, Y) can be easily generated using the Box–Muller algorithm (see Example 2.2.2). ||

We can implement this example with the following code. It is worth noting that I try to keep the same form between this two programming language as much as possible. For example, Julia does not support arbitrary mean and variance in its Gaussian sampling function `randn()`, while R can directly realize in `rnorm()`, but we both adopt the linear transformation of Gaussian distribution.

R

```
bigibbs <- function(T, rho)
{
  x = numeric(T+1)
  y = numeric(T+1)
  for (t in 1:T){
    x[t+1] = rnorm(1) * sqrt(1-rho^2) + rho*y[t]
    y[t+1] = rnorm(1) * sqrt(1-rho^2) + rho*x[t+1]
  }
  return(list(x=x, y=y))
}
## example
res = bigibbs(2e6, 0.5)
```

Julia

```
function bigibbs(T::Int64, rho::Float64)
  x = ones(T+1)
  y = ones(T+1)
  for t = 1:T
    x[t+1] = randn() * sqrt(1-rho^2) + rho*y[t]
    y[t+1] = randn() * sqrt(1-rho^2) + rho*x[t+1]
  end
  return x, y
end

## example
x, y = bigibbs(Int64(2e6), 0.5)
```

Results

The running environment is as follows:

- System: Ubuntu 18.04 (Windows subsystem for Linux)
- Processor: Intel(R) Core(TM) i7-6700 CPU @ 3.40GHz x 8
- Memory: 16 GiB
- R version: 3.4.4 (2018-03-15)
- Julia version: 1.0.0 (2018-08-08)

In terminal, use `time` command to get their running time:

```
weiya $ time julia toy_gibbs.jl

real    0m0.410s
user    0m0.422s
sys     0m0.406s
weiya $ time Rscript toy_gibbs.R

real    0m6.506s
user    0m6.406s
sys     0m0.094s
```

Obviously, in our toy example, Julia outperforms much than R, nearly 16 times. Try another number of iterations, the results are similar.

Moreover, we can use `PyPlot` to plot in Julia v1.0, such as the histogram in this toy example:

